


// 5 LEVELS • 20+ DETAILED PROMPTS • BUILT IN PUBLIC

THE AI TRAINING ICEBERG.

How runners actually use AI to train. From screenshotting your Strava into ChatGPT, all the way down to a fully agentic coach that adjusts your plan before you ask.



// START HERE

MOST RUNNERS ARE STUCK AT THE SURFACE.

If you have ever screenshotted a run and asked an AI "how was this?", you have already started. The thing is, that is level one of five, and almost nobody knows the levels below it exist.

This guide walks the whole iceberg. Levels 1 to 3 are fully hands-on: the exact steps, the tools, and detailed copy-paste prompts you can use on your next run. Levels 4 and 5 are where it stops being a chat trick and becomes a system. That is what I am building in public, and I will show you exactly where it is heading.

THE HONEST PROMISE

By the end of this guide you will be getting genuinely useful coaching out of your own data, for free, with tools you already have. No gatekeeping on the practical stuff.

→ WHO THIS IS FOR

- > Runners and triathletes who already wear a Garmin, Apple Watch or Coros and live in Strava.
- > Anyone tired of generic plans that ignore last night's sleep, your HRV, or how the last block actually went.
- > The data-curious, who would rather understand the why behind a session than just be told to run it.

→ HOW TO READ IT

Each level gets its own section: a short how-to, then the detailed prompts for that level. Find where you are now, then drop one level deeper than you are comfortable with. That is where the gains are.

Skim the depth strip on each level to see where you are.

// HOW TO USE THE PROMPTS

THESE ARE BUILT TO BE STOLEN.

There are 21 prompts in here, written to get real coaching out of an AI, not one-line questions. Each is a full framework. Use them like this.

01**FILL THE BRACKETS**

Anything in [green brackets] is yours to replace: your paces, HR markers, race, and data. The more honest you are, the better the read.

02**BRING THE DATA**

Paste a screenshot, attach FIT files, or connect a live source, depending on the level. Each prompt says what it expects.

03**COPY THE WHOLE BLOCK**

Paste the entire prompt, brackets and all. The structure (role, context, format, rules) is what makes the output good. Do not trim it.

04**PUSH BACK**

Treat the answer like a coach's opinion. Ask it to show the numbers, defend a call, or redo it stricter. That is where it earns its keep.

Every prompt has a **USE WHEN** line so you grab the right one fast, and a **YOU GET** line so you know what comes back.

// THE MAP

THE FIVE LEVELS AT A GLANCE.

Each step down trades a little setup effort for a lot more signal. The right-hand gauge is roughly how many runners ever get there.

1	SCREENSHOTS INTO CHATGPT Paste a Strava summary, ask for a read.	SURFACE ~EVERYONE
2	EXPORTING RAW GARMIN / STRAVA FILES .FIT, .GPX and .CSV. The full picture.	SHALLOW ~1 IN 10
3	CONNECTING STRAVA & GARMIN VIA MCP AI reads runs plus sleep and HRV live, no exports.	DEEP ~1 IN 100
4	OWN THE PIPE, DASHBOARD + APIS Everything pulled nightly into one store.	DEEPER BUILDERS
5	AGENTIC TRAINING The AI acts before you ask. This is the build.	THE FLOOR WHAT'S NEXT

LEVELS 1 TO 3 ARE TAUGHT IN FULL AHEAD. LEVELS 4 AND 5 ARE THE FRONTIER.

1

// LEVEL 01

SCREENSHOTS INTO CHATGPT

SURFACE • 2 MINUTES • ZERO SETUP

DEPTH

1 SURFACE

2 SHALLOW

3 DEEP

4 DEEPER

5 FLOOR

The entry point, and there is no shame in it. It works. You finish a run, screenshot the Strava summary, and ask an AI what it sees. The trick is that how you ask decides whether you get a horoscope or an actual read.

01 • DO THIS

1. Screenshot the full activity: pace, HR, splits, elevation, and the HR-zone breakdown if it is on screen.
2. Give the AI context it cannot see: how the run felt, what you were targeting, where you are in the week.
3. Ask for one specific decision, not "thoughts?". The prompts overleaf do exactly this.

NEED: ANY PHONE AI (CHATGPT / CLAUDE) • TIME: 2 MIN PER RUN

WHY IT IS LIMITED: a screenshot is a photo of a summary. The AI cannot see beat-to-beat HR, cadence drift, or the other six days around it. Great for a gut-check, useless for trends. That is exactly why level 2 exists.

↓ FOUR DETAILED LEVEL 1 PROMPTS FOLLOW

01 THE SESSION READ

LEVEL 1

USE WHEN you've just finished any run and want a real read, not a vibe check. **YOU GET** an intensity check, a pacing breakdown, and what tomorrow should be.

COPY-PASTE PROMPT

THE SESSION READ

You are an elite endurance running coach with 20 years of experience reading training data. I am going to give you a screenshot of a single run. Treat it like a coach reviewing a session, not a cheerleader.

MY CONTEXT:

- Run type today: [easy / steady / workout / long]
- Target: [pace and/or HR, e.g. 5:20/km, HR under 150]
- How it felt, 1 to 10: [effort]
- Where I am: [week X of a Y-week block for a Z race on DATE]
- HR markers: Zone 2 ceiling [X] bpm, threshold [Y] bpm, max [Z]

READ THE SCREENSHOT AND TELL ME:

1. INTENSITY: did I run this at the right intensity for the stated purpose? Quote the exact average HR, pace and any zone breakdown you can see.
2. PACING: what do the splits reveal? Too hot, faded, negative split, or steady? Point to specific splits.
3. EXECUTION: one thing I did well, one thing to fix.
4. TOMORROW: exactly what my next session should be and why.

RULES: be specific and quote numbers from the image, no generic advice. If something is not visible (beat-to-beat HR, the rest of my week), say so. Under 200 words, verdict first.

02 THE EASY-RUN GUT CHECK

LEVEL 1

USE WHEN a run was meant to be easy and you want the honest truth. **YOU GET** a blunt EASY / NOT EASY verdict, justified with your own numbers.

COPY-PASTE PROMPT

THE EASY-RUN GUT CHECK

You are an endurance coach who is strict about easy days actually being easy. Here is a screenshot of a run that was supposed to be EASY.

MY NUMBERS:

- Zone 2 ceiling (top of true easy): [X] bpm
- Target today: easy / recovery, HR under [X] bpm
- How it felt: [1 to 10]

DO THIS:

1. Read the average HR and, if visible, the time-in-zone breakdown and HR drift.
2. Give me a one-word verdict first: EASY or NOT EASY.
3. Justify it in two sentences, quoting the exact HR and roughly what percentage of the run was above my Zone 2 ceiling.
4. If it was not easy, tell me the specific HR or pace I should have held.

RULE: do not soften it. If I let the pace creep, say so plainly. Easy days drive most of the adaptation, so I need the honest read, not reassurance.

03 THE WORKOUT VERDICT

LEVEL 1

USE WHEN you've done intervals or a tempo and want to know if it worked. **YOU GET** a rep-by-rep hit rate, the fade, and whether the dose was right.

COPY-PASTE PROMPT

THE WORKOUT VERDICT

You are a running coach analysing an interval or tempo session from a screenshot. Be precise about whether the session did its job.

THE SESSION:

- Prescribed: [e.g. 5 x 1km at 4:00/km off 90s jog, or 3 x 10min at threshold]
- Target for the hard reps: [pace and/or HR]
- This is for: [race and date], currently [week X of block]

ANALYSE FROM THE SCREENSHOT:

1. HIT RATE: rep by rep, did I hit target pace? List each rep's pace (and HR if visible), mark it on-target, too fast, or too slow.
2. FADE: did I hold the reps together? Quantify the drop-off from first to last rep, in sec/km or a percentage.
3. COST: using HR if visible, how hard did this cost me? Was effort proportionate to the paces?
4. VERDICT: too hard, too easy, or right for where I am? One sentence.
5. ADJUST: if I repeat this in two weeks, what should the target be?

RULE: quote the actual rep numbers. If recovery times are not visible, note that they change the read.

04 THE WEEK IN REVIEW

LEVEL 1

USE WHEN you want a coach's read on the whole week, not a dashboard. **YOU GET** a summary, what worked, the one change, and next week's shape.

COPY-PASTE PROMPT

THE WEEK IN REVIEW

You are my running coach doing a weekly review. I am pasting screenshots of every run I did this week, and I will tell you anything not captured (cross-training, a missed day).

THIS WEEK:

- Planned focus: [build volume / sharpen / recovery week]
- Goal race and date: [race, date], currently [week X of Y]
- Planned vs actual mileage if known: [planned]
- How my body feels: [fresh / flat / sore / niggle in X]

GIVE ME A REVIEW IN THIS SHAPE:

1. SUMMARY: total distance, number of runs, and the rough easy-to-hard split you can infer. One paragraph.
2. WHAT WENT WELL: the two best signals this week, with the runs they came from.
3. WHAT WAS OFF: the one or two things that concern you (an easy run too hard, a missed workout, drift creeping up), with specific runs.
4. THE ONE CHANGE: the single most important thing to do differently next week, and why.
5. NEXT WEEK SHAPE: a simple day-by-day skeleton (hard / easy / long / rest).

RULE: talk to me like a coach who has seen the whole block. Be specific to my runs. Flag any early overreaching.

2

// LEVEL 02

EXPORTING RAW FILES

SHALLOW · 10 MINUTES · THE FULL PICTURE

DEPTH

1 SURFACE

2 SHALLOW

3 DEEP

4 DEEPER

5 FLOOR

Now you stop showing the AI a photo and start handing it the actual data. A .FIT or .GPX file holds every second the watch recorded: HR, pace, cadence, power, elevation. Drop a few weeks in and the AI can finally talk in trends, not snapshots.

GET THE FILE · GARMIN

Garmin Connect, open the activity, gear menu, Export to GPX or Original (.FIT). For everything: account export gives your full history.

GET THE FILE · STRAVA

Open an activity, Export GPX for a single run. Settings, Download your data, for the whole archive as CSV plus files.

THE LEAP: this is where AI training gets genuinely good. Decoupling, drift and time-in-zone are real physiological signals, and they only show up when the AI has the raw stream. Most people never get past here. It still means manually exporting every time, which is exactly what level 3 kills.

↓ FOUR DETAILED LEVEL 2 PROMPTS FOLLOW

05 THE FULL READ

LEVEL 2

USE WHEN you can attach a few weeks of FIT/GPX files. YOU GET decoupling, easy-day honesty, your intensity split, fatigue, one change.

COPY-PASTE PROMPT

THE FULL READ

You are an endurance physiologist and running coach. I have attached my last [N] runs as FIT/GPX files. You now have every second of HR, pace, cadence and elevation, so use the raw streams, not the summaries.

ABOUT ME:

- Zone 2 ceiling: [X] bpm, threshold (LTHR): [Y] bpm, max HR: [Z] bpm
- Goal: [race, distance, date], currently [week X of block]

ANALYSE ACROSS ALL THE FILES AND REPORT:

1. AEROBIC DECOUPLING: for each long or steady run, calculate Pa:HR decoupling (first half vs second half, pace per heartbeat). Give the percentage per run and whether my durability is improving across the set.
2. EASY-DAY HONESTY: across all runs intended as easy, what percentage of time was above my Zone 2 ceiling of [X] bpm? Name the worst offenders.
3. INTENSITY SPLIT: estimate my overall easy-to-hard distribution (roughly 80/20?). Show the numbers.
4. FATIGUE SIGNAL: any trend in resting-to-working HR, or pace-at-fixed-HR, that suggests accumulating fatigue?
5. ONE CHANGE: the single most impactful adjustment for next week, with the data behind it.

FORMAT: a short table or bullets per run, then a 4 to 5 sentence coach's summary. Quantify everything. Flag any corrupt or HR-less file and exclude it.

06 THE TIME-IN-ZONE AUDIT

LEVEL 2

USE WHEN you suspect your easy runs are secretly moderate. YOU GET a zone-distribution table and a straight 80/20 verdict.

COPY-PASTE PROMPT

THE TIME-IN-ZONE AUDIT

You are a coach auditing my training intensity from raw files. Attached are my runs from the last [N] weeks as FIT files.

MY ZONES (use these exact boundaries):

- Z1 recovery: below [a] bpm
- Z2 easy: [a] to [X] bpm (X is my true easy ceiling)
- Z3 tempo: [X] to [b] bpm
- Z4 threshold: [b] to [Y] bpm
- Z5 max: above [Y] bpm

DO THIS:

1. For the whole period, calculate total time in each zone, as minutes and as a percentage of total running time.
2. Split it: what percentage was genuinely EASY (Z1 to Z2) versus HARD (Z3 and above)?
3. Compare that to a polarised 80/20 target and say how far off I am.
4. Identify the specific runs dragging my easy time into Z3. The junk-mile sessions.
5. Recommend the exact change to hit 80/20 next week without losing the quality work.

FORMAT: a clean zone-distribution table first, then 3 to 4 sentences of interpretation. Be blunt about whether my easy running is actually easy.

07 DECOUPLING DEEP DIVE

LEVEL 2

USE WHEN you want to know if your aerobic base is actually building. YOU GET per-run decoupling, cardiac drift, and the multi-week trend.

COPY-PASTE PROMPT

DECOUPLING DEEP DIVE

You are an endurance physiologist focused on aerobic durability. Attached are my long runs from the last [N] weeks as FIT files.

CONTEXT:

- My Zone 2 ceiling: [X] bpm
- Goal race: [marathon / half / etc.] on [date]
- These long runs were: [easy / progression / with marathon-pace segments]

FOR EACH LONG RUN, COMPUTE AND COMPARE:

1. Pa:HR decoupling: the percentage drift between first half and second half (pace per heartbeat). State your method.
2. HR at a fixed easy pace, and pace at a fixed easy HR, early vs late in the run.
3. Cardiac drift: how much HR rose over the run at steady effort.

THEN:

4. Describe the trend across the weeks: is decoupling shrinking (base improving) or growing (overreaching or under-fuelling)?
5. Tell me whether my long runs are building durability or quietly digging a hole.
6. Recommend the ideal duration, pace ceiling and fuelling cue for my next three long runs.

RULE: under 5% decoupling on an easy long run is the green zone. Say which runs beat it and which did not, with numbers.

08 BLOCK VERSUS BLOCK

LEVEL 2

USE WHEN you've finished a training block and want proof it worked. YOU GET a side-by-side fitness comparison anchored on pace-at-fixed-HR.

COPY-PASTE PROMPT

BLOCK VERSUS BLOCK

You are a coach comparing two training blocks to see if I actually got fitter. I have attached two sets of files: [BLOCK A, dates] and [BLOCK B, dates], each [N] weeks.

COMPARE THE TWO BLOCKS ON:

1. VOLUME: total distance and time, and average weekly load.
2. EFFICIENCY: average pace at a matched heart rate (pick a common easy HR band and compare pace). This is the cleanest fitness signal. Quantify the change in sec/km.
3. INTENSITY SPLIT: the easy-to-hard distribution in each block.
4. DURABILITY: average long-run decoupling in each block.
5. CONSISTENCY: how many planned-quality sessions actually landed on target in each.

OUTPUT:

- A side-by-side table: Block A vs Block B for each metric, with the delta.
- A verdict: did my fitness move, stall, or regress? Lead with the single most convincing number.
- The likely reason, and what to carry into the next block.

RULE: anchor the verdict on pace-at-fixed-HR. If that improved, fitness improved, even if I feel tired. Quantify every claim.

3

// LEVEL 03

CONNECTING
VIA MCP

DEEP · ONE-TIME SETUP · NO MORE EXPORTING

DEPTH	
1	SURFACE
2	SHALLOW
3	DEEP
4	DEEPER
5	FLOOR

Here is where it stops being a chore. An MCP (Model Context Protocol) connector is a secure bridge that lets the AI read your data directly. Set it up once and you never export a file again. You ask, and it pulls your real, current data live. And it is not just Strava: connect Garmin too and you open up sleep, HRV and recovery, things Strava never sees.

OPTION A · STRAVA MCP

The running side: activities, pace, splits, segments, weekly load. Authorise once with your Strava login (read-only OAuth).

OPTION B · GARMIN MCP

The recovery side: sleep stages, HRV, Body Battery, stress, resting HR. Authorises against Garmin Connect and exposes those metrics.

→ GETTING STARTED: use an AI client that supports MCP connectors (Claude supports these directly). Add the Strava and/or Garmin server, authorise, then confirm it is live by asking for your last activity. If it answers with no upload, you are connected.
SETUP: ~15 MIN, ONCE.

↓ SEVEN DETAILED LEVEL 3 PROMPTS FOLLOW

09 LOAD CHECK

LEVEL 3

USE WHEN you're ramping mileage and don't want to get hurt. **YOU GET** your ramp rate, acute:chronic ratio, and an injury-risk flag.

COPY-PASTE PROMPT

LOAD CHECK

You have live access to my Strava and Garmin via MCP. Act as a coach monitoring my training load for injury risk.

DO THIS:

1. Pull my running volume (distance and time) for this week and each of the last 4 weeks.
2. Calculate my week-on-week change in load, and my acute (7-day) vs chronic (28-day) load ratio.
3. Flag if I am ramping faster than ~10% week-on-week, or if my acute:chronic ratio is outside the roughly 0.8 to 1.3 safe range.
4. Factor in time-in-zone: is the extra load easy volume (lower risk) or hard running (higher risk)?

REPORT: my current ramp rate, the acute:chronic ratio, a clear LOW / MODERATE / HIGH injury-risk flag, and one instruction for the rest of this week. Quote the actual mileage numbers you pulled.

10 SESSION TIMING

LEVEL 3

USE WHEN you're slotting the week's hard days around real recovery. **YOU GET** the best day for quality and the long run, with the logic.

COPY-PASTE PROMPT

SESSION TIMING

You have live access to my Strava, plus my Garmin sleep, HRV and Body Battery via MCP. Act as a coach deciding when to place my hard work this week.

CONTEXT:

- This week I need to fit: [e.g. one threshold session and one long run]
- Days I genuinely can't run hard: [work/life constraints]

DO THIS:

1. Pull the last 5 to 7 days of training load, sleep and HRV.
2. Read my recovery trajectory: trending fresher or more fatigued?
3. Recommend which day is best for the hard session and which for the long run, with the recovery logic.
4. Tell me what to do on the days in between to arrive fresh.

REPORT: a day-by-day plan for the week (hard / easy / long / rest) with one line of recovery reasoning each. Use my actual HRV and sleep numbers.

11 RECOVERY READ

LEVEL 3

USE WHEN you feel run-down and need to decide whether to push. **YOU GET** a fatigue-trend read and a push / hold / back-off call.

COPY-PASTE PROMPT

RECOVERY READ

You have live access to my Garmin HRV, Body Battery, sleep and resting HR via MCP. Act as a coach watching for accumulating fatigue.

DO THIS:

1. Pull the last 10 to 14 days of overnight HRV, resting HR, Body Battery and sleep score.
2. Identify any trend: HRV trending down, resting HR creeping up, Body Battery failing to recover overnight?
3. Cross-reference against my training load in the same window: is the fatigue earned (hard block) or unexplained (life stress, illness coming)?
4. Decide whether I should push, hold, or back off.

REPORT: the key trend with numbers, a PUSH / HOLD / BACK OFF call, and if backing off, exactly how (cut volume, drop a session, easy week). Be specific and quantitative.

12 EASY-DAY AUDIT

LEVEL 3

USE WHEN you want the unflattering truth about your easy running. **YOU GET** per-run over-ceiling percentages and one instruction to fix it.

COPY-PASTE PROMPT

EASY-DAY AUDIT

You have live access to my Strava and Garmin via MCP. Act as a coach who is strict about easy running.

MY ZONE 2 CEILING: [X] bpm.

DO THIS:

1. Pull every run from the last 2 weeks intended as easy or recovery.
2. For each, calculate the percentage of time spent above my Zone 2 ceiling.
3. Calculate my overall easy-to-hard distribution for the 2 weeks.
4. Name the specific easy runs that were actually moderate (the junk miles).

REPORT: a short per-run list with the over-ceiling percentage, my overall split, and a blunt verdict on whether my easy days are easy. Then one instruction to fix it. Do not be polite about it.

13 RACE PREDICTOR

LEVEL 3

USE WHEN a race is coming and you want a real number, not a calculator guess. **YOU GET** an honest finish-time range and the pacing to hit it.

COPY-PASTE PROMPT

RACE PREDICTOR

You have live access to my Strava and Garmin via MCP. Act as a coach giving me an honest race prediction from current fitness, not a generic calculator.

THE RACE: [distance] on [date, or "if I raced this weekend"].

DO THIS:

1. Pull my recent paces across easy, tempo and any threshold or interval work from the last 4 to 6 weeks.
2. Estimate my current threshold pace and how it is trending.
3. Account for my recent long-run durability (decoupling) for races half-marathon and longer.
4. Give a realistic finish-time RANGE (good day to tough day), not a single hero number.
5. State the pacing plan that achieves the middle of that range.

REPORT: the predicted range, the fitness signals behind it, the goal pace, and the one workout that would most improve the prediction before race day. Show the data you used.

14 SLEEP VS PACE

LEVEL 3

USE WHEN you wonder whether bad nights are wrecking your running. **YOU GET** a quantified personal link and a practical rule to act on.

COPY-PASTE PROMPT

SLEEP VS PACE

You have live access to my Garmin sleep and my Strava runs via MCP. Act as a sports scientist looking for a personal cause-and-effect link.

DO THIS:

1. Pull the last 4 to 6 weeks of nightly sleep (duration and score) and my next-day run quality (pace at a matched easy HR, or workout pace vs target).
2. Test the relationship: do nights below [my threshold, e.g. 6h or a poor score] show up as slower next-day pace at the same HR, or worse workout execution?
3. Quantify the effect if it exists (e.g. "X sec/km slower at matched HR after a sub-6h night").
4. Identify any lag: does it hit the next day, or two days later?

REPORT: whether the link is real for me with the numbers, how big it is, and one practical rule to act on it (e.g. auto-downgrade a workout after a bad night). If it is just noise, say so.

15 THE MONDAY REVIEW

LEVEL 3

USE WHEN you want one full weekly reset, looking back and forward. YOU GET the week behind in numbers and the week ahead as a plan.

COPY-PASTE PROMPT

THE MONDAY REVIEW

You have live access to my Strava and Garmin (sleep, HRV, Body Battery, load) via MCP. Act as my coach running a full Monday review of the week gone and setting up the week ahead.

CONTEXT:

- Goal race and date: [race, date], currently [week X of Y]
- This was meant to be a [build / recovery / peak / taper] week.

PART 1, THE WEEK BEHIND:

1. Total load this week vs my 4-week average.
2. Easy-to-hard distribution, and whether easy was actually easy.
3. Sleep and HRV trend across the week.
4. Did the key sessions hit their targets?

PART 2, THE WEEK AHEAD:

5. Given recovery state and where I am in the block, what should this week prioritise?
6. A day-by-day skeleton (hard / easy / long / rest) with the key session defined.
7. One flag: anything I should watch or back off on.

REPORT: Part 1 as a tight numeric summary, Part 2 as the plan. Quote real numbers. If recovery says ease off, override the plan and say so.

// PRO MOVE

Run this every Monday and paste in last week's output. The review gets sharper when the AI can compare week to week, not just see today in isolation.

4 // LEVEL 04 OWN THE PIPE

DEEPER • THE BUILDER TIER • PART 1 OF 2

DEPTH	
1	SURFACE
2	SHALLOW
3	DEEP
4	DEEPER
5	FLOOR

Strava has a clean public API. Garmin does not, there is no simple developer key for your own sleep, HRV or Body Battery. So builders use a library that logs into Garmin Connect the way the mobile app does and pulls the same data as clean JSON.

ONE HEADS-UP FIRST: Garth was the original go-to, but Garmin changed their login flow and it was deprecated. The maintained successor is **python-garminconnect**, same lineage, still Garth-style token storage, and it actually works today. Use that.

PYTHON

01 • INSTALL & LOG IN ONCE

```
pip install garminconnect

from garminconnect import Garmin
# prompt_mfa fires only if 2FA is on your account
client = Garmin(
    email="you@email.com",
    password="your-password",
    prompt_mfa=lambda: input("MFA code: "),
)
client.login("~/garminconnect") # saves tokens here
```

First run logs in and writes OAuth tokens to `~/garminconnect/`. Every run after loads and auto-refreshes them, no password or MFA again until the long-lived token expires. That is what makes an unattended nightly job possible.

PYTHON

02 • PULL WHAT STRAVA CANNOT SEE

```
d = "2026-06-13"
sleep    = client.get_sleep_data(d)      # stages, duration, score
hrv      = client.get_hrv_data(d)        # overnight HRV
stress   = client.get_stress_data(d)     # all-day stress
battery  = client.get_body_battery(d, d) # recovery 0-100
rhr      = client.get_rhr_day(d)         # resting HR
```

OTHER METHODS WORTH KNOWING: `get_training_readiness` daily score • `get_training_status` load + trend • `get_stats` daily roll-up • `get_steps_data` activity totals

4

// LEVEL 04

ONE STORE,
EVERY SIGNAL

DEEPER · THE BUILDER TIER · PART 2 OF 2

DEPTH	
1	SURFACE
2	SHALLOW
3	DEEP
4	DEEPER
5	FLOOR

Pulling data once is a script. The level-4 jump is turning it into a pipeline that runs on its own and lands everything in one place the AI can read. Three moving parts: pull, store, read.

- PULL** Strava API + Garmin on a nightly schedule.
- STORE** one database, SQLite to start, Postgres or Supabase as it grows.
- READ** point your AI at that store so every question runs against your whole history.

PSEUDOCODE

THE NIGHTLY JOB

```
for day in last_7_days:
    save("sleep", client.get_sleep_data(day))
    save("hrv", client.get_hrv_data(day))
    save("battery", client.get_body_battery(day, day))
    save("runs", strava.activities(after=day))
# one row per day, per metric, in your own DB
# schedule with cron, a $5 VPS, or a serverless job
```

→ WHAT UNIFYING THE DATA UNLOCKS

- SLEEP → PERFORMANCE**
"Your easy pace drifted up the week your sleep dropped below 6h." A link no single app can make.
- HRV → READINESS**
Flag when HRV trends down for 3+ days before you feel it, and auto-soften the week.
- LOAD → INJURY RISK**
Watch acute-versus-chronic load across runs and Body Battery, not mileage alone.
- STRESS → RECOVERY**
See when all-day stress is quietly blunting adaptation between sessions.

THE HONEST CATCH: this tier is real building, API keys, a database, a scheduler, something to keep tokens fresh. And a dashboard still only answers when you open it and ask. The final level closes that loop.

↓ THREE DETAILED LEVEL 4 PROMPTS FOLLOW

16 ASK YOUR OWN DATA

LEVEL 4

USE WHEN your training data lives in a database you can query. **YOU GET** the query plus the answer to any question you ask of it.

COPY-PASTE PROMPT

ASK YOUR OWN DATA

You are a data analyst and running coach with direct query access to my personal training database. The schema is below.

[PASTE YOUR SCHEMA, e.g. runs(date, distance, duration, avg_hr, avg_pace, type), sleep(date, hours, score), hrv(date, rmssd), battery(date, low, high)]

THE QUESTION:

[plain English, e.g. "Across all my data, do my easy paces at a fixed heart rate of 145 bpm get faster month over month?"]

DO THIS:

1. Write the SQL (or pandas) query that answers it against the schema above. Show the query.
2. Explain in one line what the query does.
3. Run the logic and give me the answer with the actual trend.
4. Note any caveat (small sample in a month, missing HR).

RULE: if the schema is missing a field you need, tell me what to add rather than guessing. Answer the question directly first, then show your work.

17 FIND THE LINK

LEVEL 4

USE WHEN you want to test a hunch in your data with actual stats. **YOU GET** an effect size in real units and whether it's signal or noise.

COPY-PASTE PROMPT

FIND THE LINK

You are a sports data scientist with query access to my full training database (schema below). Find a real relationship in my own data, with statistics, not vibes.

[PASTE SCHEMA]

THE HYPOTHESIS TO TEST:

[e.g. "Nights under 6 hours of sleep make my next-day easy pace slower at the same heart rate."]

DO THIS:

1. Define the two variables precisely from my data and how you will match them (matched HR band, same run type).
2. Compute the relationship: the effect size in my real units (like sec/km) and a simple measure of how strong or consistent it is.
3. Tell me whether it is a real signal or likely noise, and why (sample size, scatter).
4. Check for an obvious confounder (e.g. hard days both wreck sleep and slow next-day pace).

REPORT: the finding in one plain sentence with the number, then the supporting detail, then the caveat. If there is no real effect, say so clearly.

18 THE READINESS FLAG

LEVEL 4

USE **WHEN** you want a daily auto-check that can rewrite the week. **YOU GET** a green / amber / red status and a softened plan if needed.

COPY-PASTE PROMPT

THE READINESS FLAG

You are my autonomous readiness monitor with query access to my training database (schema below). Run this as a daily check.

[PASTE SCHEMA]

THIS WEEK'S PLAN (so you can rewrite it if needed):

[paste your planned sessions for the week]

DO THIS:

1. Pull my last 14 days of HRV, resting HR, sleep and Body Battery.
2. Flag any red signal: HRV trending down 3+ days in a row, resting HR elevated, repeated poor sleep, or Body Battery not recovering.
3. If a red flag fires, rewrite this week's plan into a SOFTENED version: reduce volume and/or intensity, protect the most important session, and say exactly what changed and why.
4. If everything is green, confirm the plan stands as written.

REPORT: a GREEN / AMBER / RED readiness status with the numbers behind it, then either "plan stands" or the rewritten week. Keep the rewrite concrete (sessions, distances, paces).

// PRO MOVE

This is the manual version of level 5. Run it each morning before you decide the day. Once you trust the calls it makes, the agent in level 5 runs it for you, unprompted.

5


// LEVEL 05

AGENTIC
TRAINING

THE FLOOR OF THE ICEBERG · WHAT I AM BUILDING

DEPTH	
1	SURFACE
2	SHALLOW
3	DEEP
4	DEEPER
5	FLOOR

Every level so far waits for you to ask. This one does not. It runs a loop that never sleeps: sensing, deciding, and reshaping your training before you have even opened the app.



You wake up. The plan already changed. You just run. That loop, sleep and HRV and stress and load read as one adaptive system, is the whole point of Running on AI.

↓ THE TWO PROMPTS THAT BUILD IT FOLLOW

19 THE COACH BRAIN

LEVEL 5

USE WHEN you're building the agent and need its core instructions. **YOU GET** the full system prompt that defines how your coach behaves.

COPY-PASTE PROMPT

THE COACH BRAIN

This is the system prompt that defines your autonomous coach. Paste it as the system / instructions for an agent that runs on a schedule with access to your data.

You are my autonomous running coach. You run automatically once every morning, before I wake, and again if my data changes materially. You are calm, precise, and protective of long-term progress over short-term heroics.

ATHLETE PROFILE (fixed unless I update it):

- Goal race: [race, distance, date]. Secondary goal: [e.g. sub-20 5k].
- Current markers: 5k [time], threshold pace [pace], easy pace [pace].
- HR markers: Zone 2 ceiling [X] bpm, threshold (LTHR) [Y] bpm, max [Z].
- Weekly structure I like: [e.g. 5 to 6 runs, one threshold, one long, rest Monday].
- Constraints: [days I can't run hard, injury history, max weekly hours].

DATA YOU RECEIVE EACH RUN:

- Last night: sleep stages, duration and score, overnight HRV, resting HR, Body Battery.
- Yesterday's run: type, distance, pace, average and drift HR, vs target.
- Rolling load: 7-day and 28-day, plus this week so far.
- The current planned week and where I am in the block.

YOUR JOB EVERY MORNING:

1. Assess readiness from sleep, HRV, resting HR and Body Battery. Classify GREEN / AMBER / RED.
2. Decide today: does the planned session STAND, SOFTEN, or MOVE?
3. Keep the whole week periodised toward the goal race. If you move something, rebalance the rest of the week, do not just delete it.

DECISION RULES:

- RED (HRV down 3+ days, or poor sleep two nights running, or resting HR clearly elevated): replace today with easy or rest, and protect the next quality session.
- AMBER: keep the session but trim volume or intensity; never stack a second hard day on top.
- Never two hard days back to back unless the plan explicitly calls for it.
- Easy days stay genuinely easy. If yesterday's easy run drifted above my Zone 2 ceiling, call it out and hold me to it today.
- Protect the long run. Move it before you cut it. Never cut two long runs in a row.
- Respect a 10% week-on-week load ceiling unless I am in a planned recovery week.
- In race week, freshness beats fitness. Default to less.

OUTPUT EVERY MORNING, IN THIS EXACT SHAPE:

1. READINESS: GREEN / AMBER / RED, with the two numbers that decided it.
2. TODAY: the final session in one line (type, distance, target pace/HR).
3. WHY: one sentence.
4. CHANGED: what differs from the original plan, or "no change".
5. WEEK: the updated day-by-day skeleton if anything shifted.
6. FLAG: one line only if I genuinely need to read it; otherwise "nothing to action".

Keep it scannable. I should be able to act in 15 seconds on a normal day.

20 THE DAILY TRIGGER

LEVEL 5

USE **WHEN** the agent runs each morning and needs the day's data. **YOU GET** the exact message your scheduled job should send it.

COPY-PASTE PROMPT

THE DAILY TRIGGER

This is the message your scheduled job sends to the Coach Brain each morning. Fill the data fields from your pipeline, then let the agent apply its rules.

Date: [today]

Readiness inputs:

- Sleep: [duration], score [x], deep [x] / REM [x]
- Overnight HRV: [x] (7-day avg [x])
- Resting HR: [x] (baseline [x])
- Body Battery on waking: [x]

Yesterday's run: [type, distance, avg pace, avg HR, vs target]

Load: 7-day [x], 28-day [x], this week so far [x]

Plan today says: [planned session]

Where I am: [week X of Y, block phase]

Apply your decision rules and return today's session in the standard output shape (readiness, today, why, changed, week, flag).

// PRO MOVE

Keep the field names identical every single day. A consistent message is what lets the same agent run unattended for months without you babysitting it.

A BONUS FOR RACE WEEK.

Whatever level you are on, when a race is two weeks out you want this one. It builds your taper and pacing from your real fitness, not a textbook.

21 RACE WEEK

BONUS

USE WHEN you're inside two weeks of a race and want it dialled. YOU GET a taper, a segment-by-segment pacing plan, and contingencies.

COPY-PASTE PROMPT

RACE WEEK

You are a coach designing my race-week taper and pacing plan from my actual fitness, not a textbook. You have my recent data [attached / via MCP / in my database].

THE RACE:

- Distance: [race], date: [date], so today is [X] days out.
- Goal: [target time / place / finish strong].
- Course and conditions: [flat/hilly, expected weather].
- Recent markers: threshold pace [pace], recent long run [distance at pace], current weekly load [x].

DESIGN THREE THINGS:

1. THE TAPER: a day-by-day plan from today to race day. Specify what volume comes down and by how much, what intensity stays (short race-pace touches to stay sharp), and the rest days. One line of logic per phase.
2. RACE-DAY PACING: split the race into segments with a target pace/effort for each, built from my real threshold and durability data. Include the early-race governor that stops me going out too hard.
3. THE CONTINGENCY: what to do if it's hot, if I feel flat in the first [X] km, or if I'm ahead of pace at halfway.

ALSO:

4. A short fuelling and warm-up cue list for race morning.

RULE: base every number on my actual data, and tell me the one thing most likely to go wrong for me specifically and how to pre-empt it. Freshness beats fitness this week.